
TO	Reader (value-investor lens)
FROM	LLM Council
RE	McKinsey Quarterly Q1 2026, <i>Deploying Agentic AI</i> — read alongside Forrester and Gartner coverage
DATE	2026-05-04

McKinsey on AI agents — what a value investor should make of it

Bottom line up front

Read this article as a sell-side note from a consulting firm, not as independent research. McKinsey, Gartner, and Forrester all sell the same product — paid frameworks and advisory services to executives — so all three have a structural incentive to publish "you need a complex new program" pieces during a hype phase. The article's *threat list* is mostly old IT-security problems with new vocabulary. Its *prescription* is a multi-year governance program that benefits the firms selling it. But the article's *implicit forecast* — that enterprise AI-agent deployment will be gated by liability, audit, and governance layers — is probably right, and is the most investable takeaway. It implies the picks-and-shovels (identity, observability, audit, cyber insurance) get paid before the AI labs do, and that the bull-case timeline for agent-driven productivity gains is too aggressive.

How to read this article structurally

McKinsey, Gartner, and Forrester are not journalism. They are subscription-and-engagement businesses whose unit economics depend on executives believing that complex new technology requires their frameworks to navigate.

- **Gartner**¹ historically frames new tech on a Hype Cycle: peak of expectations → trough → plateau. By every visible signal — analyst calendar, vendor positioning, conference programming — agentic AI is at or just past the peak phase. The McKinsey article fits the peak-phase template: large value claim, novel risk taxonomy, framework prescription.
- **Forrester**² tends to publish more practitioner-oriented research and is generally more skeptical of pre-deployment governance overhead than McKinsey. A Forrester read of the same topic is more likely to emphasize "start small, instrument, iterate" than "build a multi-year governance stack first."

- **McKinsey** sits between the two and produces the most consultative output — large strategic claims paired with prescriptions that map naturally onto multi-month engagements.

For an investor, this means: don't read McKinsey for *whether* something is true. Read it for *what large enterprises are about to start spending money on*. The article's checklist is a reasonable preview of the procurement line items that will appear in F500 budgets over the next 12–24 months.

What the article gets right (and what's investable from it)

1. Liability is shifting to the deployer.

This is the part the article almost surfaces but doesn't name clearly. Cyber-insurance carriers are already excluding or repricing losses from autonomous AI actions. The EU AI Act, and sectoral regulators in finance and healthcare, treat the company that *deploys* an agent as the responsible party. Forrester has been more explicit about this than McKinsey. This is a real, durable change — not narrative — and it gates enterprise rollout.

Investable implications:

- Cyber insurers re-rating premiums and exclusions favors carriers with disciplined underwriting and hurts ones writing AI risk loosely.
- "AI audit / model assurance" becomes a real budget line; established big-four audit franchises and a small number of pure-plays will absorb most of it.
- Identity and access management vendors benefit because every agent needs a credential and every credential needs governance — but this is well-priced into the leaders.

2. The enterprise AI rollout is slower than the consensus narrative.

McKinsey's nine-workstream prescription — even if you reject the prescription — is a fair preview of how F500 procurement will actually behave. Each workstream is a 12–24 month effort. That means the bull case on agentic-AI productivity gains within 2–3 years is overstated for the half of GDP that is regulated and large. Forrester's research consistently shows enterprise tech adoption lags vendor narratives by 2–3 years; this cycle will not be the exception.

Investable implications:

- Revenue ramps for AI infra and model providers depending on regulated-enterprise pull-through are likely to disappoint relative to consensus on the 2-year view.

- The spending that does flow goes disproportionately to compliance-adjacent vendors, not to model labs.
- Any private-market valuation that assumes "enterprise will adopt agentic workflows in 12–18 months at scale" is using McKinsey's upside number without McKinsey's procurement reality.

3. There are two genuinely new risks the article underplays.

- **The AI's instructions can come from anything it reads.** Inputs from emails, documents, or web pages can shape behavior the same way source code does.
- **Many agents share the same model.** When a single AI model powers thousands of agent instances, one bad update or one clever attack causes correlated, simultaneous failure across all of them — closer to a crop disease or a shared-assumption banking failure than to traditional IT risk.

Investable implications: this favors providers offering monitoring, evaluation, and red-teaming of model behavior; it disfavors organizations concentrated on a single model provider without contingency. The "vendor concentration" risk at hyperscaler-tied AI shops is underappreciated.

What the article gets wrong (and what to discount)

The "\$2.6T to \$4.4T in annual value" projection.

This is a McKinsey Global Institute–style total-addressable-market figure. These projections are not investable inputs. The historical hit rate on consultant-firm 5-year TAM projections in technology is poor; the range itself (\$2.6T to \$4.4T is a 70% spread) admits as much. Treat it the way you would treat a sell-side analyst's "blue sky" scenario — directionally interesting, quantitatively useless. It is the same pattern as Gartner's mid-2010s blockchain TAM and IoT TAM projections, both of which under-delivered on the published timeline by an order of magnitude.

The "80 percent of organizations encountered risky behaviors" statistic.

A survey statistic with no published threshold for what counts as "risky." For comparison, ~100% of organizations have encountered risky behaviors from spreadsheets and email. The number is being used to justify spending on a specific framework. It does not indicate the frequency of material loss events. There is no failure case study with quantified loss anywhere in the article — a tell.

The framing "trust must be the foundation."

This is the opposite of how mature security has worked since the 1970s. Real security practice assumes things will fail and bounds the damage — *zero trust*³, *least privilege*⁴, blast-radius containment. A company that builds its agent program on a "trust foundation" frame will under-invest in the controls that actually matter (scoped credentials, reversible actions, human approval on high-impact operations) and over-invest in policy documents.

The four "novel" risks (out of five) are old problems renamed.

- *Chained vulnerabilities* = standard microservice-trust failures.
- *Cross-agent task escalation* = the "confused deputy" problem, named in 1988.
- *Synthetic-identity risk* = unauthenticated-actor problems solved by signed requests.
- *Untraceable data leakage* = missing audit logs.
- *Data corruption propagation* = bad data in, bad data out.

For an investor, this matters because the relevant capability already exists in the market. The companies that solve these problems are the established security/identity/observability franchises, not greenfield "AI governance" startups built around a new vocabulary.

Triangulating against Forrester and Gartner

Side-by-side, the value-investor signal is:

SIGNAL	MCKINSEY	GARTNER	FORRESTER	READ
Risk framework is novel	Strong	Moderate	Weak	"Novel taxonomy" framing benefits McKinsey's advisory practice.
Pre-deployment governance is mandatory	Strong	Moderate (per Hype Cycle methodology)	Weak (start-small bias)	Pure governance-first prescription is McKinsey's house style; treat it as marketing of their service line.
Liability and audit will gate rollout	Implicit	Explicit	Explicit	Cross-firm consensus and the most investable observation.
AI agents will produce \$2.6–4.4T/yr	Yes	Lower (Gartner's own AI TAM is more bounded)	Lower still	Discount this number heavily.

SIGNAL	MCKINSEY	GARTNER	FORRESTER	READ
Enterprise rollout will be fast	Implicitly slow	Explicitly slow (5–10 yr plateau)	Explicitly slow	Consensus across all three: slower than the AI-bull narrative.

What an investor should actually do with this

1. **Don't trade the headline numbers.** The \$2.6–\$4.4T figure and the 80% statistic are decoration. Use them to gauge sentiment, not to size positions.
2. **Position around the procurement reality.** The companies that get paid first in an enterprise agent rollout are the ones that already sell identity, access management, observability, audit, and cyber-insurance. They get paid whether the agents themselves succeed or not.
3. **Discount agent-pure-play timelines.** Any thesis that requires regulated F500 customers to deploy autonomous agents at scale within 12–18 months is leaning on a bull-case adoption curve that all three firms — even McKinsey — implicitly contradict. The realistic ramp is 3–5 years for the regulated half of the economy.
4. **Track the liability and insurance signal more than the framework signal.** Cyber-carrier exclusions, EU AI Act enforcement actions, and regulator guidance on deployer liability are the real gating events. These are public, datable, and harder for analyst firms to spin.
5. **Treat "governance frameworks" as the new SOC2⁵** — a billable line item that rewards established compliance and audit franchises, mostly performatively, regardless of whether the underlying threat profile justifies it.

The McKinsey article is half consulting marketing and half a useful signal that enterprise AI-agent rollout will be gated by liability and audit — which means the durable winners over the next 3–5 years are the picks-and-shovels around identity, monitoring, audit, and cyber insurance, not the agent vendors themselves.

FOOTNOTES

1. **Gartner** — research and advisory firm. Best known for the *Magic Quadrant* (a 2x2 vendor-rating chart) and the *Hype Cycle* (a five-stage framework describing how new technologies move from inflated expectations through disillusionment to productive use). Subscription-driven; primary customer is the enterprise IT buyer.
2. **Forrester** — competing research and advisory firm. Best known for the *Forrester Wave* (a vendor-evaluation framework). Generally regarded as more practitioner- and customer-experience-oriented than Gartner, and somewhat less framework-heavy than McKinsey.
3. **Zero trust** — a security philosophy: do not assume any system, user, or program is safe just because it's inside your network. Verify every request. The opposite of "trust must be the foundation."
4. **Least privilege** — give each program (or person, or AI agent) the minimum access it needs to do its job, and nothing more. Limits damage if the program is compromised.
5. **SOC2** — a widely-adopted enterprise security audit standard. Frequently cited as an example of a control regime that is real but heavily commoditized into a checkbox-compliance industry; "the new SOC2" is shorthand for a control regime that becomes a standard line item in enterprise budgets, partly substantive, partly performative.